

# The SRE's Crystal Ball

Predicting Performance with Queues and USL



Aravindh Sampathkumar

Booking.com

October 09, 2025

# About Me

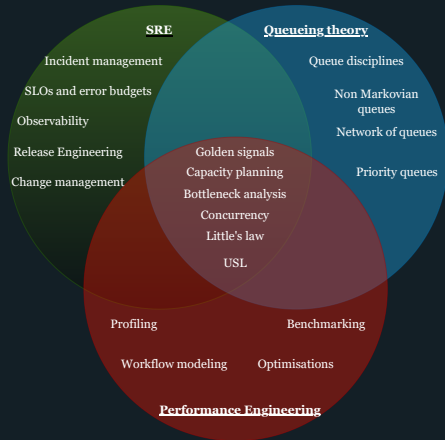
**Aravindh Sampathkumar**

*Site Reliability Engineer @ Booking.com*

## Background


- High Performance Computing (HPC)
- Storage Systems
- Performance Engineering

# The big picture



# The big picture

- Let's have some fun.
- Aimed at inspiring you, my peer practitioners to apply these concepts in ways I haven't thought of.
- A system/service is a chain of bottlenecks - removing one reveals the next.



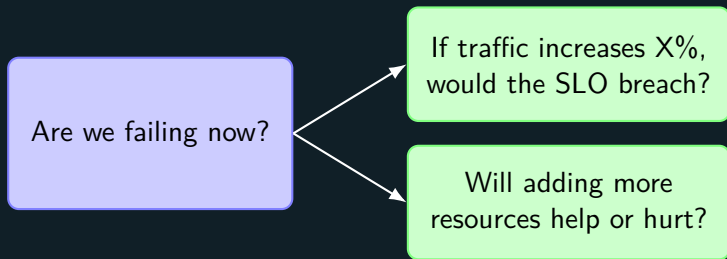
# Credits

- Dr. Neil J. Gunther authored Universal Scalability Law and taught GCAP workshop.
- Stefan Moeding maintains the R package that I use in Demo.
- Phil Larson who mentored and introduced me to applied queueing theory.

# Why bother?

Most SLOs are reactive alarms

- **Availability SLO Miss:** *"We're down."*
- **Latency SLO Miss:** *"We're slow."*



Move from firefights to preventing *some* issues entirely — think in **queues** and **scalability models**.

# Agenda

- Thinking in queues
- Golden signals through the lens of queueing theory
- Predicting scalability with Universal Scalability Law(USL)
- See it in action

# The Birth of Queueing Theory



- In the early 1900s, Danish mathematician **Agner Krarup Erlang** had to figure out how many telephone circuits were needed to handle a given number of calls without excessive waiting or dropped connections.
- His work created queueing theory — the mathematical study of waiting lines (or queues).

Figure: A. K. Erlang  
(1878-1929) Source: Wikipedia



# Thinking in queues

- Queues are everywhere - Loadbalancers, connection pooling, Jira boards..

# Thinking in queues



# Thinking in queues

- Queueing systems are non-linear - It is not intuitive. Double the instance count and you can serve double the load right?
- Queues occur even if there is enough average capacity - Grocery store
  - High variance in arrivals
  - High variance in service times

# Thinking in queues

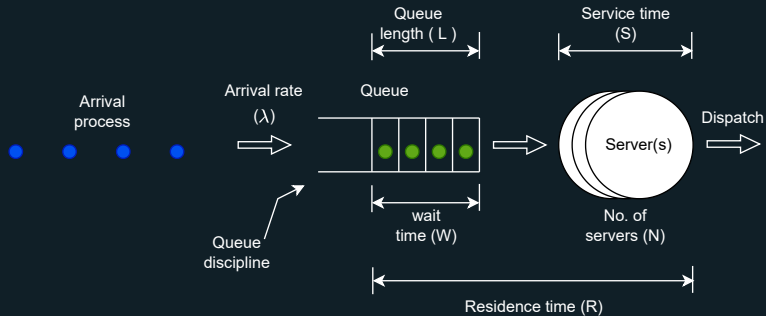


Figure: A simple queue.

# Thinking in queues - Terminology

---

Queue length ( $L$ )	Number of requests in the system or queue( $L_q$ )
Arrival Rate ( $\lambda$ )	The rate at which requests enter the system (e.g., requests/sec). aka "demand" or "load".
Service Time ( $S$ )	The time required for a single server to process one request.
Waiting Time ( $W$ )	The time a request spends waiting in the queue before its service begins.
Response Time aka Latency( $R$ )	The total time a request is in the system. The sum of waiting and service time ( $R = W + S$ ).
Utilization ( $\rho$ )	The fraction of time a server is busy.
Number of servers ( $N$ )	The number of servers in the system(node/cpu/thread/instance).
Throughput ( $X$ )	Completion rate of requests. In a stable system, ( $X = \lambda$ ).

---

# Golden signals 🕶 through queueing theory

- Latency = Service time + Waiting time
  - Utilisation drives waiting time
  - Coefficient of variation is a leading indicator
- Traffic as Arrival rate( $\lambda$ ): The demand driver
- Errors as symptoms of overload
- Saturation as high Utilisation( $\rho$ ): The harbinger of Doom

## Little's Law and Related Formulae


$$L = \lambda R$$

$$\rho = \lambda S$$

or

$$\rho = \frac{\lambda S}{N}$$


# Utilisation curve and the "Knee"






## A little intuition


Which design performs better?



## Applied Theory: The Case of the Sluggish API

 An API endpoint ('/api/v1/resource') is timing out. Latency has spiked!


# Applied Theory: The Case of the Sluggish API

 An API endpoint ('/api/v1/resource') is timing out. Latency has spiked!

What do we know?

- **Arrival Rate ( $\lambda$ ):** 180 requests/second.
- **Avg. Service Time (S):** 50 milliseconds per request.
- **Number of Servers (N):** 4 API server pods.

# Applied Theory: The Case of the Sluggish API

 An API endpoint ('/api/v1/resource') is timing out. Latency has spiked!

What do we know?

- **Arrival Rate ( $\lambda$ ):** 180 requests/second.
- **Avg. Service Time (S):** 50 milliseconds per request.
- **Number of Servers (N):** 4 API server pods.
  
- *"Should we scale up to 8 pods? "*
- *"Would it be enough?"*

## Applied Theory: The Case of the Sluggish API

$$\rho = \frac{\lambda S}{N} = \frac{\text{Arrival Rate} \times \text{Service Time}}{\text{Number of Servers}}$$
$$\rho = \frac{180 \text{ req/s} \times 0.05 \text{ s/req}}{4 \text{ pods}} = \frac{9}{4} = \mathbf{2.25}$$

## Applied Theory: The Case of the Sluggish API

$$\rho = \frac{\lambda S}{N} = \frac{\text{Arrival Rate} \times \text{Service Time}}{\text{Number of Servers}}$$
$$\rho = \frac{180 \text{ req/s} \times 0.05 \text{ s/req}}{4 \text{ pods}} = \frac{9}{4} = \mathbf{2.25}$$

- Utilisation ( $\rho > 1.0$  or 100%)  $\Rightarrow$  the system is unstable. In other words, the queue is growing infinitely.
- Increase N to 8 pods?  $\Rightarrow \rho$  will compute to 1.125 (still  $> 1.0$ ).
- Let's increase N to at least 10 pods ( $\rho = 0.9$ ) to also accommodate variance.

I promised a crystal ball. Lets do some predictions.

# USL

## ***Scalability***

A mathematical function of being able to perform more work (RPS, TPS etc) while work per server<sup>1</sup> remains constant and the number of servers increases.\*

## ***Universal Scalability Law(USL)***

A formal definition of scalability.

---

<sup>1</sup> (node/instance, cpu/thread etc)

\* Improvised from a definition by Dr. Neil J Gunther




# USL - Linear Scalability

The ideal condition. We all want that right?

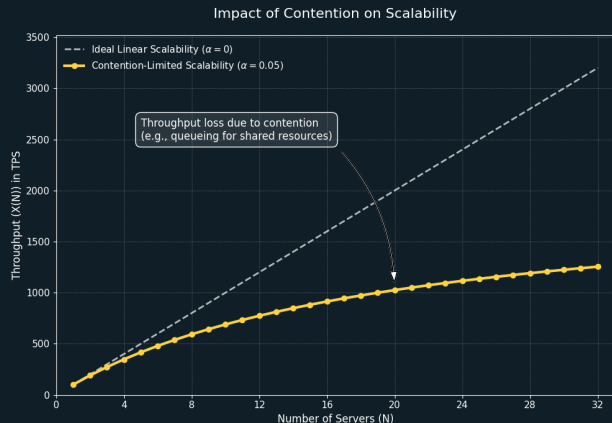
$$X(N) = \frac{\gamma N}{1}$$

Where:

- $X(N)$  is the throughput with  $N$  servers.
- $\gamma$  is the ideal throughput of a single server.
- $N$  is the number of servers.



# USL - Scalability villain no.1 - Contention




$$X(N) = \frac{\gamma N}{1 + \alpha(N - 1)}$$

Where:

- $\alpha$  represents contention.
- contention is non-parallelisable serialised work.

# USL - Scalability villain no.2 - Coherency (a.k.a. Crosstalk)



# USL

$$X(N) = \frac{\gamma N}{1 + \alpha(N - 1) + \beta N(N - 1)}$$

Where:

- $\gamma$  represents ideal single server throughput.
- $\alpha$  represents contention.
- $\beta$  represents coherency penalty.

## USL - What do we do with it?

*Model system/scalability (obviously).*

1. Apply observed system/service behaviour metrics.
2. Estimate  $\alpha$ ,  $\beta$ , and  $\gamma$  using non-linear least squares regression.
3. Interpret the results to know scalability limits and improve bottlenecks.

Demo

Source at <https://github.com/aravindhsampath/srecon25-usl>

# USL - Limits

- Noisy data - real-world data is quite noisy(network jitter, OS scheduler, GC pauses etc.).
- Distributed systems - harder to model aggregate functions(/order depends on 10 microservices).
- Asynchronous and event driven architecture pattern is hard to model.
- Coherency factor( $\beta$ ) assumes 1-1 coordination.
- Systems dont always behave the same way at scale. E.g switch to a different algorithm or co-ordinate in batches etc.
- Noisy neighbours - multi-tenancy.
- Rate limits and throttling at dependencies.

## Common pitfalls - USL

Keep these in mind while working with USL.

- Garbage In Garbage Out - be diligent about noise. Noise will have non linear effects.
- Over-extrapolation - Dont forecast too far(new constraints emerge at scale). 2x is the rule of thumb.
- Max throughput is often NOT the goal - max throughput @desired latency is.



# Key takeaways

- Start thinking in queues. Enrich your monitoring - capture service times and waiting times in histograms, queue lengths.
- Use Little's law for napkin math of fundamental relationship between Utilisation, Throughput/arrival rate, and response time.
- Use USL as a diagnostic compass to identify coarse contention and coherence penalties - rethink design choices to achieve practical goals.
- Keep queuing theory wisdom in mind while designing services:
  - Favour pooled resources - single queue and a shared pool of workers.
  - Attack variability - Use caching strategies, optimisations for slower code paths etc to reduce service time variability.
  - Heavy-tailed service times? consider priority queues or size based routing.

## References and further reading

- Dr. Neil J Gunther's [page](#) on Scalability and USL.
- Baron Schwartz's [The essential Guide to Queueing Theory](#).
- Neil J. Gunther and Stefan Moeding's [USL R package](#)
- Eben Freeman's LISA 17 talk on [Queueing theory in practice](#)
- Kavya Joshi's talk on [Applied Performance Theory](#).

# Thank You!

## Questions?

## Get in Touch

- ✉ : [aravindh at fastmail.com](mailto:aravindh@fastmail.com)
- in : [linkedin.com/in/aravindhsampath](https://linkedin.com/in/aravindhsampath)
- GitHub : [github.com/aravindhsampath](https://github.com/aravindhsampath)
- Bluesky : [aravindh.net](https://aravindh.net)

